

Fast Money, Slow Locks: Chronicles of Atomic Trespasses and Memorials of Transactional Craft

– or what is this TREM –

Otávio A Araújo Silva

@otavioarj

NullByte 8 Nov. 2025



Agenda

1 Race Conditions & Semantics

- Problem Framing
- Ledger Semantics
- Distributed Angle (Glossary)

2 Exploitation Playbook

- Attacker Model & Recon
- Concrete Exploits in Ledgers
- Failure Signals (Telemetry)
- Exploit Snippets
- Strong Writes for Money

3 Case Studies

- OpenSSH (2024) — Timing and Signals
- Starbucks (2015) — Gift-Card Double-Spend

4 Testing & Chaos

5 This never happened

6 This is the TREM

1 Race Conditions & Semantics

- Problem Framing
- Ledger Semantics
- Distributed Angle (Glossary)

2 Exploitation Playbook

3 Case Studies

4 Testing & Chaos

5 This never happened

6 This is the TREM

Race Conditions & Semantics

Problem Framing

- Latency windows in payment paths enable: double-spend, duplicate payments, stale authorizations, overdrafts.
- Parallelization (workers, retries, webhooks, autoscaling) scales both throughput *and* exploit surface.
- Monetary loss is non-linear: narrow races \times burst concurrency \Rightarrow large leakage.

Security View

- Attackers synchronize requests to the pre-commit gap (race orchestration).
- Hotspots: balance checks, hold creation, gift-card transfers, reversal/cancel flows.

Definitions: Logical vs Data Races

Logical Race - Our focus here

Correctness depends on the interleaving of operations on shared **application state** (DB rows, cache keys, queues). High level race conditions.

Data Race

Memory-model violation, unsynchronized read/write of same location causing undefined behavior. Low-level race conditions, e.g., virtual memory r/w locks violations

Why It Matters

Logical races survive **proper thread locks** and typically arise across processes/services/DBs.

Taxonomy of Race Conditions

Application/DB

- **TOCTOU**: stale check then effect (classic “check-then-act”).
- **Lost update**: concurrent writes overwrite each other; e.g. balance decrements race.
- **Write skew & phantoms**: constraints enforced via separate reads; concurrent inserts slip through.
- **ABA-like logical patterns**: value toggles and returns; naive guards pass.

Distributed System

- **Ordering races**: out-of-order queue/webhook deliveries observe inconsistent states.
- **Duplicate delivery races**: retries reapply effects if idempotency is weak.
- **Lease/lock races**: expired holders still write without fencing tokens.

Mobilization vs Settlement (Ledger Invariants)

Terminology

- **Mobilize (Hold/Authorization)**: reserve funds temporarily.
- **Settle (Capture/Transfer)**: finalize and post movement; release or consume hold.

Critical Invariants

- Same unit of value cannot be mobilized twice before settlement.
- Hold creation must not exceed available balance after concurrent holds.
- Settlement must relate to a single, valid, unconsumed hold.

Operational Concurrency: GC, Skew, Split Brain

Runtime/Infra

- **Garbage Collection pauses:** stop-the-world intervals that delay threads and expand timing windows.
- **Clock skew:** node clocks drift; time-based leases/expirations misfire (stale holders act as valid).
- **Split brain:** partition elects multiple leaders; conflicting writes occur concurrently.

Consistency Building Blocks

- **Isolation levels:** RC/RR/Serializable; weaker levels admit lost updates, write skew, phantoms.
- **Idempotency:** retries return same logical result; requires dedup state at write.
- **Exactly-once:** rarely achievable across networks; enforce invariants at the DB boundary.

Consistency Building Blocks

Isolation Levels: RC/RR/Serializable

- **RC (Read Committed)**: Transactions only read committed data; prevents dirty reads, allows non-repeatable reads and phantom reads, e.g. multiples insert/deletes inconsistency.
- **RR (Repeatable Read)**: Guarantees that if a transaction reads a value, subsequent reads will return the same value; prevents non-repeatable reads, still allow phantom reads.
- **Serializable**: Strongest isolation level; transactions execute as if they were run one after another; prevents all anomalies: lost updates, write skew, and phantoms.
- **Weaker levels trade-off**: Lower isolation levels= performance++. Admit anomalies: concurrent overwrites), write skew (constraint violations), and phantoms, e.g new rows appearing in range queries).

Consistency Building Blocks

Idempotency

- **Definition:** An operation is idempotent when executing it multiple times produces the same logical result as executing it once.
- **Retry safety:** Enables safe retries in distributed systems where network failures may cause duplicate requests; critical for at-least-once delivery semantics.
- **Deduplication state:** Requires maintaining state to detect and filter duplicate operations; typically implemented using request IDs, timestamps, or sequence numbers stored at write time.
- **Examples:** SET operations are naturally idempotent; INCREMENT operations require deduplication tokens to become idempotent.

Consistency Building Blocks

Exactly-Once Semantics

- **The challenge:** True exactly-once delivery is theoretically impossible across networks due to fundamental distributed systems problems (network partitions, timeouts, crashes).
- **Practical approach:** Instead of guaranteeing exactly-once at the network level, enforce invariants and consistency rules at the database boundary.
- **DB-level guarantees:** Use transactions, unique constraints, and conditional writes to ensure logical exactly-once semantics even with at-least-once delivery.
- **Implementation pattern:** Combine idempotency keys with transactional boundaries; the database becomes the source of truth for deduplication and consistency enforcement.

1 Race Conditions & Semantics

2 Exploitation Playbook

- Attacker Model & Recon
- Concrete Exploits in Ledgers
- Failure Signals (Telemetry)
- Exploit Snippets
- Strong Writes for Money

3 Case Studies

4 Testing & Chaos

5 This never happened

6 This is the TREM

How Attackers Find and Time Races

Recon

- Identify multi-step flows: check → authorize/hold → capture/settle → release.
- Probe endpoints that move money or create ledger effects (transfer, reload, reversal).
- Observe retries/timeouts/webhook behavior; infer idempotency gaps.

Timing Orchestration

- Burst concurrent requests; jitter control to land between check and commit.
- Exploit server/client retries; force ambiguous failures (timeouts) to trigger replays.
- Abuse cross-service latencies (API→queue→worker→DB) to desynchronize guards.

Exploit: Double Mobilization (Two Holds, One Balance)

Mechanics

- Two workers **A** and **B** read “balance $\geq N$ ” before any mutation.
- Both insert holds based on the same stale snapshot; weak isolation (RC/RR) admits it.
- Synthetic liquidity emerges; settlement later drains real funds or creates debt.

Failure Conditions

- Split check/mutate across calls/services.
- No single-statement guard; no unique idempotency key tied to effect.

How Races Surface in Metrics

Quantitative Indicators

- Rising rate of **duplicate-key conflicts** (idempotency table) vs throughput.
- Ratio: holds created vs holds released/captured in short intervals.
- Inconsistent ledger deltas (balance) across replicas (local copies); spike in p99 latency^a at transactions endpoints.

^aHighest latency of 1% slowest response

Qualitative Indicators

- Chargeback-like symptoms without external processor issues.
- Intermittent “insufficient funds” after successful authorizations.

Anti-Pattern (Pseudo-SQL) — Split Check and Mutate

Don't: Predicate Not Bound to Effect

```
-- Workers A and B execute concurrently
bal = SELECT balance FROM accounts WHERE id = :u;
IF bal >= :n THEN
    INSERT INTO holds(user, amount) VALUES(:u, :n);
END IF;
```

Where It Fails (A & B)

- **A** and **B** both observe the same pre-state **before** any mutation.
- The check is detached from the write; interleavings admit **two holds**.
- Under RC/RR isolation, no serialization prevents the lost-update style outcome.

Atomic Predicate + Effect (SQL Pattern)

Do: Bind Predicate to Effect

```
BEGIN;  
UPDATE accounts  
SET balance = balance - :n  
WHERE id = :u AND balance >= :n;  
-- rowcount == 1 -> success; else -> insufficient funds  
INSERT INTO holds(user, amount, idempotency_key)  
VALUES(:u, :n, :k);  
COMMIT;
```

Why This Works

- Single-statement guard; success gated by affected row count: **no separate read**.
- Effect coupled to a unique **idempotency key**; replays = same result.

Idempotency & Async Events

Transactional Idempotency

- Unique key per logical funds request; persisted with the effect **in the same transaction**.
- On replay: return stored result; never reapply state changes.

Treat Async Events as Participants

- Webhooks/timeouts fenced/serialized with the write (version checks, fencing tokens).
- Reject mid-flight state flips unless tied to the same commit boundary.

1 Race Conditions & Semantics

2 Exploitation Playbook

3 Case Studies

- OpenSSH (2024) — Timing and Signals
- Starbucks (2015) — Gift-Card Double-Spend

4 Testing & Chaos

5 This never happened

6 This is the TREM

OpenSSH (2024): Signal-Handler Race — CVE-2024-6387 “regr~~e~~SSHion”

Essence

- Timing-sensitive interaction between signal delivery and handler logic in code: `LoginGraceTime(handler) → async-signal-safe func(ex.,syslog) → malloc/free` = RCE
- Asynchronous events flip flags mid-operation; assumptions diverge pre/post handler.
- Analogy: expirations/timeouts/webhooks behave like signals; can preempt money writes.

Portable Lessons

- Treat async events as preemption points; design critical sections robust to interruption.
- Avoid “flag-based” correctness outside the transaction boundary.

Starbucks Gift-Card Double-Spend (2015): Mechanics

Observed Behavior

- Concurrent **reload/transfer** calls allowed duplicate credit between gift cards.
- Guard and effect split across services (app → API → DB) admitted interleavings.
- Idempotency not enforced end-to-end; replays landed as new effects.

Attacker Steps (Reconstructed)

- ① Pick two cards: source (A) and destination (B). Ensure A passes the guard.
- ② Fire k parallel transfers/reloads A→B with identical parameters.
- ③ Induce retries via client timeouts or network jitter to widen the replay surface.
- ④ Observe B credited multiple times while A debits are inconsistent/singular.

Starbucks (2015): Root Causes & Exploit Envelope

Root Causes

- Non-atomic *check* (balance on A) and *mutate* (debit A, credit B).
- Weak/implicit idempotency (no unique key or dedup record tied to effect).
- Missing single-writer invariant for the pair (A,B) during transfer.

Exploit Envelope

- Window width: sub-10 ms sufficient under burst; amplified by autoscaling workers.
- Reliable under client-induced retries and partial timeouts.
- Detector: transient mismatches of #transfers vs net ledger delta per minute.

1 Race Conditions & Semantics

2 Exploitation Playbook

3 Case Studies

4 Testing & Chaos

5 This never happened

6 This is the TREM

Temporal Fuzzing & Chaos for Money Paths

Technique

- Inject jitter between guard and effect; N workers hammer the same logical request.
- Expectation: zero duplicate holds/captures; any non-zero indicates a race.
- Metrics: dedup hit rate, conflict retries, holds vs releases/captures ratio.

Chaos Scenarios

- Kill a replica mid-transaction; drop/reorder webhooks; expire sessions mid-write.
- Force client retries/timeouts; verify invariants under maximal interleaving.

This never happened?



This is The TREM



```
Conn: [REDACTED]:443
Matched k1: ae169a71-9695-46aa-b9b9-72964f35813a

Conn: [REDACTED]:443
HTTP 200
Matched k1: 0f5494e9-bb1d-42aa-935b-f4a6340682d0

Conn: [REDACTED]:443
HTTP 200
HTTP 200
Matched k1: 17fb1d74-e87b-40c8-83a8-f78ac48890eb

Conn: [REDACTED]:443
Matched k1: 17fb1d74-e87b-40c8-83a8-f78ac48890eb

Conn: [REDACTED]:443
HTTP 201
HTTP 502
```

TREM - Transactional Racing Executor Monkey



Transactional Racing Executor Monkey v0.8 Pu

```
-d int
    Delay ms between reqs.
-h string
    Host:port override; default is addr from Host HTTP Header.
-k
    Keep-alive connections; persist TLS tunnels for every request, including while looping (-xt N).
-l string
    Comma-separated request RAW HTTP/1.1 files.
-mode string
    Mode: sync or async. (default "async")
-o
    Save *Last* (per-thread) HTTP response as out_<timestamp>_t<threadID>.txt
-px string
    HTTP proxy; http://ip:port
-re string
    Regex (Golang) definitions file. Each line applies to a request file, respectively.
    Examples:
        One Regex line for each request file, e.g., line 1 will use regexes for request 1
        Format: regex1':key1$$regex2':key2$$...regex':keyN supports multiples regex per line/request
        Note: Use backtick (`) character, not (').
-th int
    Threads count. (default 1)
-u string
    Universal replace key=val every request; e.g., !treco!=Val replaces !treco! to Val in every request, multiple times if matched.
-x int
    When looping, chain request x to N, where x is -l [1..x..N], default disabled.
    Ex: -l "req1,req2,req4" -x 2, does reqs 1 to 4, then iterates -xt times from req2 to req4 (default -1)
-xt int
    Requests loop count:
        0=infinite
        -1= zero loops
```

Questions?

Jack the baboon worked on the railway system in South Africa for 9 years without ever making a single mistake.



This is TREM



otavioarj/TREM